How to sell AI front desk
https://youtu.be/mAcsHilWPjc?si=ovkEcjrZTFnfjka7

Product overview

Name: DemoDrop AI (working title)
Type: Multi-tenant SaaS for agencies to auto-generate AI chat and voice demo agents for local businesses in seconds.

Core idea:
Given a business URL (and minimal metadata), the system automatically:

- Scrapes the website
- Builds a structured knowledge base
- Generates AI chat + voice agents tailored to that business
- Publishes a branded demo page
- Delivers the demo link via email/SMS/WhatsApp
- Tracks opens, interactions, and pipeline status for the agency

---

Key value proposition

- For agencies: A repeatable “10-second demo” engine to sell AI agents at recurring fees (e.g. \$500–\$1,000/month).
- For businesses: A concrete, personalized preview of how AI captures missed calls, chats, and after-hours leads—without any setup on their side.

Feature set

1. Lead intake & targeting

- Lead capture forms:
  - Inputs: Business name, website URL, contact name, email, phone, niche/vertical.
- Bulk import:
  - CSV upload or CRM sync to import many businesses at once.
- AI prospect scoring (optional v2):
  - Detects if site already has chat/voice; flags “no AI” prospects.

2. Website scraping & knowledge base

- Crawler:
  - Fetches main pages (home, services, contact, FAQ, pricing, about).
- Content extraction:

  - Cleans HTML, removes boilerplate, extracts structured entities (services, locations, hours, phone numbers, CTAs).
- Knowledge base builder:
  - Stores chunks in a vector database with metadata (page type, importance, language).
  - Supports multi-language content if detected.

3. AI agent generation

- Chat agent:
  - System prompt template with pluggable variables: niche, tone, brand voice, business goals.
  - Uses KB + tools (e.g. “check opening hours”, “capture lead details”).
- Voice agent:
  - Same core logic as chat, wrapped with:
    - STT (speech-to-text)
    - TTS (text-to-speech) with configurable voice, language, and speaking style
    - Call flow rules (greeting, fallback, escalation, handoff to human).
- Agent configuration UI:
  - Parameters: tone, aggressiveness in booking, lead qualification questions, escalation rules, languages, allowed actions.
  - Vertical presets: dental, real estate, restaurants, clinics, etc.

4. Demo page generation

- Branded landing page per business:
  - Business logo, colors (auto-extracted from site where possible).
  - Embedded chat widget (web) and click-to-call demo (voice).
  - Short explainer copy: “Here’s how much revenue you’re missing when you don’t answer calls/chats.”
- Lead capture on demo page:
  - Captures business owner’s reactions, questions, and “ready to install” signals.

5. Outreach & delivery

- Automated demo delivery:
  - Email/SMS/WhatsApp templates with personalized copy and unique demo link.
- Sequences:
  - Follow-up messages if demo not opened after X hours/days.
  - Branching: opened but no reply, interacted heavily, requested call, etc.

6. Tracking & CRM pipeline

- Event tracking:
  - Demo page opens, time on page, chat messages, voice calls, form submissions.
- Pipeline stages:
  - New lead → Demo sent → Demo opened → Engaged → Proposal sent → Closed won/lost.
- Dashboard:

  - Per-agency and per-client metrics: open rates, engagement, conversion to paid, MRR.
- Integrations:
  - Webhooks + API for syncing with external CRMs (including GoHighLevel, HubSpot, etc.).

7. Subscription & billing

- Agency accounts:
  - Multi-tenant: each agency has its own workspace, branding, and client list.
- Pricing models:
  - Per-seat, per-agent, or usage-based (number of demos, minutes of voice, etc.).
- White-label options:
  - Custom domain, logo, colors for agencies.

Technical specification (for your AI developer)

1. Architecture

- Frontend:
  - SPA (React/Vue/Svelte) for agency dashboard and configuration.
  - Public demo pages rendered as lightweight, fast-loading pages.
- Backend (API layer):
  - REST/GraphQL for:
    - Lead management
    - Scraping jobs
    - KB management
    - Agent configuration
    - Demo page generation
    - Analytics & events
- Background workers:
  - Queue-based system for:
    - Scraping & parsing
    - KB embedding
    - Demo page generation
    - Outreach sending
    - Analytics aggregation

2. Core services

- Scraper service:
  - Respect robots.txt and rate limits.
  - Configurable depth and page limits.
- NLP/Embedding service:
  - Chunking, embedding, and storing in vector DB.
- LLM orchestration:
  - Prompt templates for:
    - KB structuring

  - Agent persona generation
  - Vertical-specific behaviors
- Chat runtime:
  - Retrieval-augmented generation (RAG) over KB.
  - Tools for lead capture, booking, escalation.
- Voice runtime:
  - Telephony integration (e.g. SIP/VoIP provider) or web-based voice widget.
  - STT + TTS pipeline with low latency.

3. Data model (high level)

- Agency
- User (within agency)
- BusinessLead (prospect)
- WebsiteSnapshot (raw HTML + parsed content)
- KnowledgeBaseItem (chunk + embedding + metadata)
- AgentConfig (chat + voice settings)
- DemoPage (URL, template, status)
- OutreachMessage (channel, status, timestamps)
- Event (open, click, chat message, call, form submission)
- Deal/PipelineStage

System flow diagram (text version)

You can paste this into a diagramming tool (e.g. draw.io, Excalidraw, Miro) as a guide.

`text
```text
[Agency User]
    |
    v
[Create Lead: Business Name + URL + Contact]
    |
    v
[Scraper Service]
    |
    v
[Website Snapshot]
    |
    v
[KB Builder: Chunk + Embed + Store]
    |
    v
[Knowledge Base for Business]
    |
    v
[Agent Generator]
```

```
   |----------------------------|
   v                            v
[Chat Agent Config]        [Voice Agent Config]
   |                            |
   |                            v
   |                      [Voice Runtime]
   v
[Demo Page Generator]
   |
   v
[Personalized Demo Page URL]
   |
   v
[Outreach Engine (Email/SMS/WhatsApp)]
   |
   v
[Business Owner Receives Demo Link]
   |
   v
[Opens Demo Page]
   |
   v
[Interacts with Chat/Voice Agent]
   |
   v
[Event Tracking Service]
   |
   v
[CRM/Pipeline Dashboard]
`
```

“Better than the original” enhancements

If you want to leapfrog the GoHighLevel-style approach, here are some differentiators you can bake in from day one:

- Open-source friendly core:
  - Pluggable LLMs, vector DBs, and telephony providers so agencies can self-host or choose vendors.
- Privacy & compliance controls:
  - Data residency options, PII redaction, configurable retention.
- Vertical playbooks:
  - Pre-built flows for top niches in your markets (e.g. Lagos-based clinics, schools, real estate, restaurants).
- Multi-channel demos:
  - Web chat, WhatsApp, Instagram DMs, and voice in one unified demo.

- “Install in 1 click” path:
  - From demo to production: one button to generate embed code, WhatsApp integration instructions, or call-forwarding setup.

If you tell me which stack your developer prefers (Node/Python, specific LLM providers, telephony options), I can turn this into a more concrete technical blueprint with endpoint examples and data schemas.

For a Front Desk Voice Assistant, the most critical factor is latency. If there is more than a 1-second delay between the user finishing a sentence and the assistant responding, the conversation feels unnatural and "robotic."
If I were building this today for a production environment (like a hotel, dental office, or office building), this is the stack I would use:

1. The Language: Node.js (TypeScript)

While Python is the king of AI research, Node.js is significantly better for a real-time voice assistant.

- Why: Voice assistants rely heavily on WebSockets and asynchronous audio streaming. Node's event-driven architecture handles high-concurrency, real-time audio buffers more efficiently and with lower overhead than Python’s typical sync/async patterns.

2. The Telephony: Vapi or Retell AI

Don't build on raw Twilio unless you have a dedicated team for WebRTC engineering.

- Recommendation: Vapi or Retell AI.
- Why: They handle the "plumbing" (noise cancellation, echo, interruption handling, and SIP/Phone trunking). They allow you to focus on the "Brain" while they ensure the audio stream from the phone is crystal clear and low-latency.

3. The "Brain" (LLM): GPT-4o Realtime

- Recommendation: OpenAI GPT-4o Realtime API.
- Why: Until recently, you had to chain 3 separate models (STT -> LLM -> TTS). GPT-4o Realtime combines them into one native audio-to-audio model. This removes the 500ms–1s delay caused by converting text back and forth.
- Alternative for Cost: If the Realtime API is too expensive, use Groq (Llama 3.1). It is the fastest text-based LLM on the market, which compensates for the time lost in the STT/TTS steps.

4. Speech-to-Text (STT): Deepgram

- Recommendation: Deepgram Nova-2.
- Why: It is currently the "gold standard" for speed and accuracy. It can handle heavy accents and noisy backgrounds (common for front desks) and returns text in under 200ms.

5. Text-to-Speech (TTS): Cartesia or ElevenLabs

- Recommendation: Cartesia (for speed) or ElevenLabs (for emotion).
- Why:

- Cartesia is built for extreme low-latency (under 100ms). Perfect for a front desk where speed is everything.
- ElevenLabs sounds the most human but is slightly slower.

🏗️ Summary Architecture

Component	Provider
Orchestration	Node.js / TypeScript
Telephony	Vapi (or Retell)
STT	Deepgram
Logic	GPT-4o (Realtime)
TTS	Cartesia (Sonic)
Database	Supabase (PostgreSQL)
Interruption Handling	Vapi's Built-in Smart Endpointing

The "Pro" Move:

Use a Function Calling (Tools) architecture. For a front desk, the assistant needs to *do* things (book a room, check an appointment, look up a customer). Since you're already using Supabase, your Node.js server can trigger Supabase Edge Functions to hit your DB or external APIs whenever the LLM decides it needs to "Check the calendar."
Would you like me to draft a basic architecture diagram or a sample Node.js starter for this?

……
Perfect — now that you've shared your preferred stack, I can turn your architecture summary into a clean, developer-ready blueprint that your AI engineer can immediately act on. I'll keep it structured, visual, and aligned with the components you listed.

Below is a refined architecture, a full product description, a technical specification, and a flow diagram — all tailored to:

- Node.js / TypeScript orchestration
- Vapi or Retell for telephony
- Deepgram STT
- GPT-4o Realtime for logic

- Cartesia Sonic for TTS
- Supabase (PostgreSQL) for DB
- Vapi smart endpointing for interruption handling

This is the kind of document you can hand directly to your developer.

🚀 PRODUCT DESCRIPTION (Developer-Ready)

Name: AI Demo Engine (placeholder — you can rename later)

A multi-tenant SaaS platform that allows agencies to instantly generate personalized AI chat and voice demo agents for any local business using only the business website URL. The system scrapes the website, builds a knowledge base, generates a branded demo page, and sends it to the business owner via email/SMS/WhatsApp. Agencies track engagement and convert demos into paying clients.

🧩 CORE FEATURES

1. Lead Intake
- Input: Business name, website URL, contact info, niche.
- Optional: Bulk CSV import.
- Auto-detects business type and missing info.

2. Website Scraping & Knowledge Base
- Scrapes key pages (Home, Services, About, Contact, FAQ).
- Cleans and chunks content.
- Embeds into vector DB (Supabase pgvector).
- Extracts structured data: services, hours, locations, pricing, FAQs.

3. AI Agent Generation
- Chat agent using GPT-4o Realtime + RAG.
- Voice agent using:
  - Telephony: Vapi or Retell
  - STT: Deepgram
  - TTS: Cartesia Sonic
  - Logic: GPT-4o Realtime
- Configurable tone, aggressiveness, lead qualification, escalation rules.

4. Demo Page Generation
- Auto-branded landing page with:
  - Business logo + colors (auto-extracted)
  - Chat widget
  - Click-to-call voice demo
  - Short pitch explaining value

5. Outreach Automation
- Sends demo link via:
  - Email
  - SMS
  - WhatsApp
- Follow-up sequences based on engagement.

6. Analytics & Pipeline
- Tracks:
  - Demo opens
  - Chat interactions
  - Voice calls
  - Lead submissions
- Pipeline stages:
  - New → Demo Sent → Opened → Engaged → Proposal → Closed

7. Multi-Tenant Agency System
- Each agency has:
  - Workspace
  - Branding
  - Client list
  - Billing

🏗️ SYSTEM ARCHITECTURE (Aligned With Your Stack)

1. Orchestration Layer (Node.js / TypeScript)
Handles:
- API endpoints
- Webhooks (Vapi/Retell events)
- Scraping jobs
- KB building
- Demo generation
- Outreach automation
- Analytics ingestion

Framework options:
- NestJS (recommended for modularity)
- Express (lighter, faster to build)

---

2. Telephony Layer

Option A: Vapi
- Smart endpointing (handles interruptions)

- Built-in call routing
- WebRTC + PSTN support
- Webhooks for:
  - call.started
  - call.ended
  - transcript.updated
  - user.interrupted

Option B: Retell
- Lower latency
- Better for natural conversations
- Similar webhook model

---

3. Speech Pipeline

STT: Deepgram
- Real-time streaming STT
- Low latency
- Multi-language support

TTS: Cartesia Sonic
- Ultra-fast neural TTS
- Natural prosody
- Supports emotional tone + speed control

---

4. Logic Engine

GPT-4o Realtime
Used for:
- Conversation logic
- Lead qualification
- RAG responses
- Interrupt handling
- Call flow decisions
- Chat agent behavior

---

5. Database Layer

Supabase (PostgreSQL)
Tables:
- agencies
- users

- leads
- website_snapshots
- kb_chunks
- agent_configs
- demo_pages
- outreach_messages
- events
- pipeline_stages

Vector search:
- pgvector extension for embeddings

---

6. Scraper + KB Builder
- Node.js scraper (Cheerio + Playwright)
- Chunking + embedding
- Store embeddings in Supabase

---

7. Demo Page Engine
- Static page generator (Next.js or plain HTML)
- Embeds:
  - Chat widget
  - Voice call widget (Vapi/Retell)
- Auto-branding from website colors/logo

---

8. Outreach Engine
- Email (Resend / SendGrid)
- SMS (Twilio)
- WhatsApp (Meta Cloud API)
- Automated sequences

🔁 SYSTEM FLOW DIAGRAM (Developer-Ready)

`
```
[Agency User]
    |
    v
[Create Lead: Business Name + URL]
    |
    v
[Scraper Service (Node.js)]
```

```
   |
   v
[Website Snapshot Stored in Supabase]
   |
   v
[KB Builder]
   |
   v
[Chunk + Embed + Store in pgvector]
   |
   v
[Agent Generator]
   |                  |
   v                  v
[Chat Agent Config]   [Voice Agent Config]
   |                  |
   |                  v
   |             [Vapi/Retell Runtime]
   v
[Demo Page Generator]
   |
   v
[Demo URL Created]
   |
   v
[Outreach Engine Sends Demo Link]
   |
   v
[Business Owner Opens Demo]
   |
   v
[Chat or Voice Interaction]
   |
   v
[Events Logged → Supabase]
   |
   v
[Agency Dashboard + Pipeline]
`
```

📦 DEVELOPER-READY SPECIFICATION

API Endpoints

POST /leads
Create new business lead.

POST /scrape
Trigger website scraping.

POST /kb/build
Chunk + embed content.

POST /agent/generate
Generate chat + voice agent configs.

POST /demo/create
Generate demo page.

POST /outreach/send
Send demo link.

POST /events/webhook
Receive:
- Vapi/Retell events
- Chat events
- Page open events

Realtime Voice Flow
1. Call starts → Vapi webhook → Node.js
2. Node streams audio → Deepgram STT
3. STT text → GPT-4o Realtime
4. GPT-4o response → Cartesia Sonic TTS
5. TTS audio → Vapi → Caller
6. Interruptions handled by Vapi endpointing
7. Events logged to Supabase

Chat Flow
1. User message → Node.js
2. Retrieve relevant KB chunks
3. GPT-4o Realtime generates response
4. Response returned to chat widget
5. Events logged

🎯 NEXT STEP
If you want, I can now generate:

A. A polished PDF-ready product brief

B. A full technical architecture diagram (boxes + arrows)

C. A developer onboarding document

D. A feature roadmap (MVP → V1 → V2)

Just tell me which one you want next.

……Got you—let's build this out in the sequence I proposed.

A. Product brief (PDF-ready narrative)

Product name
Working title: DemoDrop AI (or "AI Demo Engine")

One-liner
A multi-tenant platform that lets agencies generate personalized AI chat and voice demo agents for any local business in minutes—using just the business website URL.

Target users
- Primary: Marketing agencies, AI consultants, call-center resellers.
- Secondary: SaaS resellers, niche operators (dental, real estate, clinics, restaurants).

Core problem
Local businesses lose leads every day from missed calls, slow responses, and no after-hours coverage. Agencies know AI agents can fix this, but selling the idea is hard without a concrete, tailored demo.

Solution
DemoDrop AI turns a business website into a live, branded AI demo:

1. Agency enters a business URL and contact details.
2. System scrapes the website and builds a knowledge base.
3. It auto-generates:
  - A chat agent (web widget)
  - A voice agent (phone/WebRTC demo)
4. It creates a branded demo page for that business.
5. It sends the demo link via email/SMS/WhatsApp.
6. It tracks opens, interactions, and pipeline status so the agency can close the deal.

Key value propositions

- For agencies
  - 10-minute demo creation instead of custom builds per prospect.
  - Repeatable, scalable sales engine for AI retainers.

  - Multi-tenant dashboard to manage all clients and deals.

- For local businesses
  - See and hear “their own” AI agent using their content.
  - No setup required—just click a link and interact.
  - Clear visibility into how many leads AI can capture.

Core features

- Lead intake & targeting
  - Single lead form (business name, URL, contact, niche).
  - Bulk import via CSV.
  - Auto-classification of niche and missing info.

- Website scraping & knowledge base
  - Scrapes key pages (Home, Services, About, Contact, FAQ).
  - Cleans and chunks content.
  - Embeds into a vector store for retrieval.

- AI agent generation
  - Chat agent powered by GPT-4o Realtime + RAG.
  - Voice agent using Vapi/Retell + Deepgram STT + Cartesia Sonic TTS.
  - Configurable tone, aggressiveness, lead questions, escalation rules.
  - Vertical presets (dental, real estate, clinics, restaurants, etc.).

- Demo page generation
  - Auto-branded landing page with logo and colors.
  - Embedded chat widget and click-to-call voice demo.
  - Short explainer copy and lead capture form.

- Outreach automation
  - Sends personalized demo links via email, SMS, WhatsApp.
  - Follow-up sequences based on opens and interactions.

- Analytics & pipeline
  - Tracks demo opens, chat messages, calls, and form submissions.
  - Visual pipeline: New → Demo Sent → Opened → Engaged → Proposal → Closed.
  - Per-agency and per-client metrics.

- Multi-tenant & billing
  - Separate workspaces per agency.
  - White-label branding (logo, domain, colors).
  - Flexible pricing models (per agent, per minute, per demo).

Differentiators

- Built around real, personalized demos, not generic templates.

- Uses best-in-class stack for realtime voice (Vapi/Retell, Deepgram, Cartesia, GPT-4o).
- Designed for agencies as first-class citizens (multi-tenant, pipeline, outreach).
- Open, modular architecture—swappable providers in future.

B. Technical architecture diagram (boxes + arrows description)

You can hand this to a designer or drop it into a diagramming tool.

High-level layers

1. Client layer
   - Agency dashboard (web app)
   - Public demo pages (per business)
   - Chat widget
   - Voice call entry (phone/WebRTC via Vapi/Retell)

2. Backend orchestration (Node.js / TypeScript)
   - REST/GraphQL API
   - Webhook handlers
   - Scraper service
   - KB builder
   - Agent generator
   - Demo page generator
   - Outreach engine
   - Analytics/event collector

3. AI & voice services
   - GPT-4o Realtime (logic + RAG)
   - Deepgram (STT)
   - Cartesia Sonic (TTS)
   - Vapi or Retell (telephony/voice runtime)

4. Data & storage
   - Supabase PostgreSQL (relational data)
   - pgvector (embeddings)
   - Object storage (HTML snapshots, logs, assets)

Suggested diagram structure (text layout)

`text
```
[Agency Dashboard (Web)]          [Business Owner Browser/Phone]
        |                                  |
        v                                  v
```

```
  [API Gateway / Backend (Node.js / TS)] <---- [Demo Page + Chat Widget]
        |
        |-------------------------------|
        v                               v
 [Scraper + KB Builder]            [Outreach Engine]
        |                               |
        v                               v
 [Website Snapshots]               [Email/SMS/WhatsApp Providers]
        |
        v
 [KB Chunks + Embeddings (Supabase + pgvector)]
        |
        v
 [Agent Generator (Chat + Voice Config)]
        |                               |
        v                               v
 [Chat Runtime (GPT-4o + RAG)]      [Voice Runtime]
                                        |
                                        v
                          [Vapi/Retell Telephony Layer]
                                        |
                                        v
                          [Deepgram STT] -> [GPT-4o Realtime] -> [Cartesia Sonic TTS]
                                        |
                                        v
                                  [Caller Audio]
`
```

Core components and responsibilities

- API Gateway / Backend
  - Authentication, multi-tenant routing.
  - CRUD for agencies, leads, configs, demos.
  - Webhook endpoints for Vapi/Retell, chat events, page events.

- Scraper + KB Builder
  - Fetches and parses website HTML.
  - Extracts text, chunks, and embeds.
  - Stores in Supabase with metadata.

- Agent Generator
  - Builds system prompts and config objects for chat and voice.
  - Stores per-business agent configuration.

- Chat Runtime
  - Receives messages from chat widget.
  - Retrieves relevant KB chunks.
  - Calls GPT-4o Realtime.

  - Returns responses and logs events.

- Voice Runtime
  - Vapi/Retell handles call setup and interruption.
  - Streams audio to Deepgram STT.
  - Sends transcripts to GPT-4o Realtime.
  - Sends GPT responses to Cartesia Sonic TTS.
  - Streams audio back to caller.
  - Logs call events.

- Outreach Engine
  - Generates personalized outreach content.
  - Sends via email/SMS/WhatsApp.
  - Tracks delivery, opens, clicks.

- Analytics & Pipeline
  - Aggregates events into per-lead and per-agency views.
  - Computes funnel metrics and statuses.

C. Developer onboarding document

1. Tech stack overview

- Language: TypeScript
- Runtime: Node.js
- Framework: NestJS or Express (your choice, but NestJS recommended)
- Database: Supabase (PostgreSQL + pgvector)
- AI Logic: GPT-4o Realtime
- Telephony: Vapi or Retell
- STT: Deepgram
- TTS: Cartesia Sonic
- Frontend: Next.js or React SPA for dashboard + demo pages

---

2. Core services/modules

Backend modules (suggested):

- auth
  - Multi-tenant auth (agencies, users).
- leads
  - CRUD for leads, bulk import.
- scraper
  - Website fetch + parse.
- kb

  - Chunking, embedding, storage.
- agents
  - Chat + voice config generation.
- demos
  - Demo page creation, URL management.
- outreach
  - Email/SMS/WhatsApp sending + sequences.
- events
  - Webhook handlers, event logging.
- analytics
  - Aggregations, pipeline views.

---

3. Environment & configuration

Environment variables (examples):

- SUPABASEURL, SUPABASESERVICE_KEY
- OPENAIAPIKEY (for GPT-4o Realtime)
- DEEPGRAMAPIKEY
- CARTESIAAPIKEY
- VAPIAPIKEY or RETELLAPIKEY
- EMAILPROVIDERKEY (Resend/SendGrid)
- SMSPROVIDERKEY (Twilio)
- WHATSAPPAPIKEY (Meta Cloud API)
- APPBASEURL (for demo links and webhooks)

---

4. Data model (high-level)

- agencies: id, name, branding, billing info
- users: id, agencyid, role, email, passwordhash
- leads: id, agencyid, businessname, websiteurl, contactname, contactemail, contactphone, niche, status
- websitesnapshots: id, leadid, url, rawhtml, parsedat
- kbchunks: id, leadid, content, embedding, source_url, metadata
- agentconfigs: id, leadid, chatconfigjson, voiceconfigjson
- demopages: id, leadid, slug, url, status, created_at
- outreachmessages: id, leadid, channel, content, status, sent_at
- events: id, leadid, type, payloadjson, occurred_at
- pipelinestages: id, leadid, stage, updated_at

---

5. Key flows to implement first

1. Lead creation → scraping → KB build
2. KB → agent config generation
3. Agent config → demo page creation
4. Demo page → outreach send
5. Demo interactions → events → analytics

---

6. Local development workflow

1. Clone repo, install dependencies.
2. Run Supabase locally or connect to dev instance.
3. Run migrations to create tables.
4. Start backend server (Node.js).
5. Start frontend (Next.js/React).
6. Use mock providers or sandbox keys for Vapi/Retell, Deepgram, Cartesia.

D. Feature roadmap (MVP → V1 → V2)

MVP (first usable version)

- Lead intake
  - Single lead form (no bulk import yet).
- Scraping & KB
  - Scrape 3–5 key pages.
  - Basic chunking and embeddings.
- Chat agent
  - Simple GPT-4o Realtime + RAG chat widget.
- Voice agent
  - Basic Vapi/Retell integration with one default voice.
- Demo page
  - Simple page with chat + click-to-call.
- Outreach
  - Manual send of demo link (copy-paste URL).
- Analytics
  - Basic event logging (opens, chats, calls).

Goal: Prove end-to-end flow works and feels impressive in demos.

V1 (agency-ready)

- Lead intake
  - Bulk CSV import.
  - Niche detection and tagging.

- Scraping & KB
  - More robust scraping and error handling.
  - Better metadata (services, hours, locations).
- Agent configuration UI
  - Sliders/toggles for tone, aggressiveness, escalation.
  - Vertical presets.
- Outreach automation
  - Email + SMS sending from inside app.
  - Simple follow-up sequence (e.g., 2 reminders).
- Analytics & pipeline
  - Visual pipeline board.
  - Per-lead timeline of events.
- Multi-tenant
  - Agency workspaces with basic branding.

Goal: Agencies can run real campaigns with multiple leads.

V2 (differentiated product)

- Advanced outreach
  - WhatsApp integration.
  - Conditional sequences based on engagement.
- Multi-channel demos
  - WhatsApp and Instagram DM demo flows.
- Deeper vertical playbooks
  - Pre-built flows for top niches (dental, real estate, clinics, schools, restaurants).
- Reporting
  - Revenue impact estimates ("you missed X leads worth ≈ Y").
- Customization & white-label
  - Custom domains for demo pages.
  - Full white-label for agencies.
- Provider flexibility
  - Ability to swap STT/TTS/LLM providers in config.

If you want, next we can zoom into one slice—for example, "Voice runtime with Vapi + GPT-4o Realtime" or "KB builder design"—and turn it into a very detailed spec your developer can start coding from tomorrow. Which slice feels most critical for you right now?